



CPC357 ASSIGNMENT 2

IOT ARCHITECTURE AND SMART APPLICATIONS

SCHOOL OF COMPUTER SCIENCE

UNIVERSITI SAINS MALAYSIA

SEMESTER 1 2025/2026

ASSIGNMENT 2

Lecturer: Dr. Mohammad Nadhir Ab Wahab

Group Members:

Name	Matric No
HARINETHIRAN A/L NAGARETINAN	22300638
ARUNRAJ A/L ARMUGAM	22300686

Submission: 18th January 2026

Table of Contents

1. Introduction.....	4
2. High-Level System Architecture Diagram.....	5
Figure 1: Shows High-Level Architecture Diagram of the System	5
2.1 Hardware Layer	5
2.2 Connectivity & Communication Layer.....	6
2.3 Cloud Infrastructure Layer.....	7
2.4 Data Management Layer.....	7
2.5 Visualization Layer.....	8
3. Design Consideration.....	9
3.1. Factors Influencing Design Decisions	9
3.1.1 Sensor Selection.....	9
3.1.2 Hardware Selection.....	9
3.1.3 Scalability	9
3.1.4 Data Security.....	10
3.1.5 Connectivity & Communication	10
3.1.6 Cost Effectiveness.....	10
3.2 Compromises	11
3.2.1 Design Compromises Between Cost and Performance.....	11
3.2.2 Usability vs. Security Considerations	11
3.2.3 Real-Time Data vs. Long-Term Analysis	11
3.2.4 Flexibility vs. Specialization.....	12
4. Development Process.....	12
4.1 Setting Up Google Cloud Platform Services	12
4.1.1 Compute Engine.....	12
4.1.2 Firewall	15
4.1.3 Service Account	16
4.1.4 BigQuery.....	19
4.2 Set Up the Hardware	21
4.2.1 Arduino Code.....	21
4.3 Set Up Environment in Google Virtual Machine.....	21
4.4 Create python script & Run it	22
5. Security Consideration.....	23
5.1 Security Measure Implemented	23
5.1.1 Transport Layer Security (TLS) for MQTT Communication	23
5.1.2 Firewall Rules in Virtual Private Cloud (VPC)	23

5.1.3 Identify & Access Management (IAM)	24
5.1.4 Data Encryption in Transit & at Rest.....	24
5.1.5 Authentication & Authorization.....	24
5.2 Strategies for Data Protection & Access Control.....	25
5.2.1 Monitoring & Alerts.....	25
5.2.2 Backup & Disaster Recovery	25
5.2.3 Access Control Policies	25
5.2.4 Regular Key Rotation.....	26
5.2.5 Penetration Testing & Vulnerability	26
6. Appendix	26

1. Introduction

The **Real-Time Weather Monitoring System** is an Internet of Things (IoT)-based one which is aimed at recording and reporting on the weather conditions as they are. The system offers proper and dependable real-time information concerning vital weather conditions, namely, temperature, humidity and precipitation. This information is needed in a variety of practical uses such as agricultural planning, urban infrastructure management and disaster preparedness. The system is economical and effective because it involves **using the DHT11 sensor** of temperature and humidity and a **rain sensor** of precipitation, which will be used to monitor the weather in real-time.

The system combines these sensors and cloud services to allow **real-time transmission of data and processing**. Through this integration, it is made possible such that the data recorded by the physical sensors which will be part of weather data is relayed in a smooth manner to a centralized platform where **data can be stored and analyzed and visualized**. The scalability and reliability of the systems can be promoted with the help of the cloud-based integration and the monitoring station is able to process the growing data and to provide the same performance and accessibility.

The increasing need in real time weather details to make informed decisions is the solution to this project. The real-time availability of temperature, humidity, and rainfall could be used in the farm sector **to identify the most efficient irrigation times** to ensure the crop is not lost due to adverse weather patterns. Urban landscape In urban setup, real time weather data can **help in controlling drainage, flood control and planning of infrastructure**. Also on-time weather information is useful in emergencies and disaster planning and response situations where the prompt availability of the correct information can enhance preparedness and response efficiency.

The system emphasizes the combination of software and hardware materials so as to deliver effective and viable functionality. The hardware layer is tasked with the responsibility of sensing and gathering weather data and the software layer is the one that transmits, stores and analyses data using cloud-based services. Such a layered architecture provides a clear way of communication among system components, as well as allows expansion in the future. Other sensors or capabilities can be added with little modification to the current system.

All in all, the Real-Time Weather Monitoring System shows that IoT types of technologies can inspire effective interconnections between physical sensory devices and online platforms to provide timely and practical weather data. An effective, scalable and real-time solution to the contemporary weather surveillance, the system can provide insight into agile and proactive decision-making in various fields.

2. High-Level System Architecture Diagram

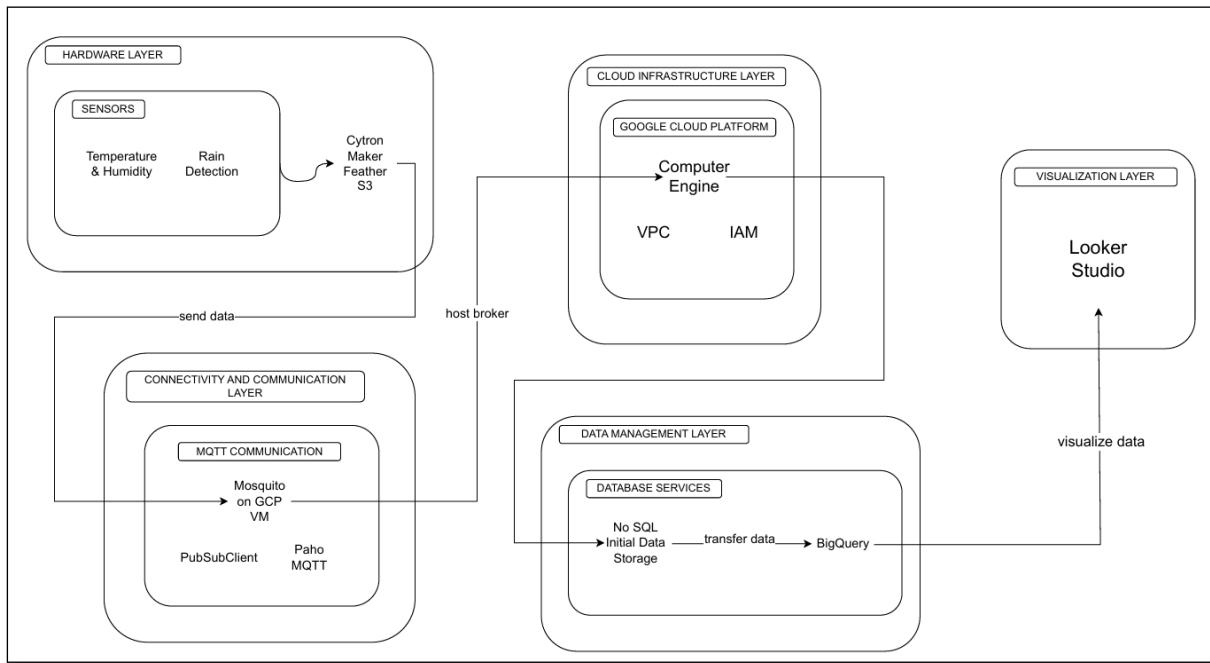


Figure 1: Shows High-Level Architecture Diagram of the System

2.1 Hardware Layer

The hardware layer is the basic part of the **Real Time Weather Monitoring System** and includes the actual IoT devices engaged in the gathering of environmental data. The main unit of this layer is a **Cytron Maker Feather S3 microcontroller** that will serve as a central unit of data collection and transmission. The two sensors fitted in this microcontroller can be used to measure the following parameters of the environment:

- **DHT11 Sensor:** This sensor is tasked with the operation of measuring the temperature and the humidity of the environment. It offers digitally and is characterized by simplicity and low cost so that it can be employed in IoT applications.

- **Rain Sensor:** This is a sensor that measures the falling of rain and sends live information about rain. It is able to predict the occurrence and the severity of rain thus the system can be used to keep track of weather conditions.

These sensors are then collated into the raw environmental data which is sent to the microcontroller which sends it to the cloud where it is further processed and analyzed. The choice of effective and reliable hardware components will provide correct data collection, and this will be the foundation on the overall functionality of the system.

2.2 Connectivity & Communication Layer

The connection and communications layer is an essential part of the Real Time Weather Monitoring System that will be in charge of the safe and efficient relay of sensor information (IoT devices) to the cloud framework. Some of the important technologies that are used in this layer include the following:

- **Mosquitto MQTT Broker:** The Mosquitto MQTT broker is deployed on a virtual machine of Google Cloud Platform (GCP) Compute Engine. It is the central messaging point, where the data between the server and IoT devices are published and subscribed to.
- **MQTT Protocol:** MQTT (Message Queuing Telemetry Transport) protocol is implemented to provide the opportunity to transmit the data with lightweight and efficiency. It is also very much applicable in IoT applications because it does not require high bandwidth. Transport Layer Security (TLS) is applied to port 8883 to achieve data integrity and security when the data is being transferred.
- **PubSubClient Library:** At the microcontroller communication, PubSubClient library is used to create MQTT communication. It allows the microcontroller to broadcast the sensor data obtained to the MQTT broker.
- **Paho MQTT Library:** The Paho MQTT library enables communication with the MQTT server on the server side to enable the system to subscribe or get incoming sensor data that can be further used.

2.3 Cloud Infrastructure Layer

The cloud infrastructure level is to be used to put up and run the critical services needed in the Real Time Weather Monitoring System. It uses a number of services provided by Google Cloud Platform (GCP) to provide scalability, security, and reliability:

- **Compute Engine:** To support data communication in real-time, a Compute Engine virtual machine which is hosting the Mosquitto MQTT broker is the high-performance, customizable computing environment.
- **Virtual Private Cloud (VPC):** VPC service is used to provide a safe and autonomous network in the system. It handles the communication of the network and implements a firewall policy, which helps to avoid the unauthorized access of MQTT broker and other cloud resources.
- **Identity and Access Management (IAM):** IAM is implemented in order to create appropriate access control and permissions in the system. BigQuery has been configured with a dedicated service account that has the privileges of the BigQuery Admin role allowing authorized management and querying of the data warehouse as well as improving the global system security.

This is used as an infrastructure layer that offers secure, scalable and reliable cloud environment that supports continuous data transmission, processing and storage of the weather monitoring solution.

2.4 Data Management Layer

A data management layer takes care of storage, processing and formatting of sensor data retrieved by the IoT devices in the Real Time Weather Monitoring System. This layer entails the following components:

- **MongoDB:** Raw sensor data are first of all saved into a database in MongoDB on a Compute engine virtual machine. Being a NoSQL database, the MongoDB database is flexible in servicing unstructured and semi-structured data. This acts as a short-lived storage measure as it allows the effective ingestion and manipulation of real-time sensor information.

- **BigQuery:** MongoDB sensor data is then sent to a data warehouse service provider, BigQuery, which is a fully managed and highly scaled service that belongs to Google Cloud Platform (GCP). BigQuery allows quick and effective querying and analysis of large volumes of data, which is why it could be applied to process and analyze a lot of data produced by the system.

This multi-layered data management strategy provides good real-time data ingestion, as well as, long-term storage and high volume analysis of weather data.

2.5 Visualization Layer

The visualization layer has a role of presenting the sensor data that has undergone the processing process in the Real Time Weather Monitoring System in a manner that is intuitive and insightful. This is done through the help of Looker Studio, an efficient business intelligence software that happens to be a product of Google Cloud Platform (GCP):

- **Interactive Dashboards:** The interactive dashboards are created with the help of Looker Studio and allow people to visualize and analyze the information related to weather in a proper manner. These dashboards provide an interactive way of searching through trends, patterns, and insights on the sensor data.
- **Integration to BigQuery:** Can be done seamlessly. Looker Studio is compatible with BigQuery, which can be accessed on the fly and has the processed information on the data warehouse. The visualizations in this integration are dynamic and automatically updated with the ingestion of new data to provide timely and correct insights.

Through the Looker Studio, the system provides actionable intelligence and a holistic outlook of the weather conditions to the benefit of improving informed decision-making by the stakeholders.

Generally, this Real Time Weather Monitoring System architecture is a well-organized and impressive IoT, which uses various services of Google Cloud Platform (GCP) to gather, transmit, store, process, and visualize the environmental data. Through a layered architecture, espoused in terms of hardware, connection and communication, cloud infrastructure, data

management, and visualization layers, the system is scalable, secure and flexible. This design allows providing useful pieces of information regarding weather, which means that it is a powerful and effective tool in real-time monitoring of the environment.

3. Design Consideration

3.1. Factors Influencing Design Decisions

3.1.1 Sensor Selection

DHT11 Sensor: DHT11 sensor was chosen because it is cheap, easy to use and has enough accuracy to measure temperature and humidity. It delivers a digital signal, which is particularly advantageous in situations with the IoT when the low cost and reliability of data obtained is the requirement. The DHT11 does not compete with more sophisticated sensors like the DHT22 in terms of accuracy; however, since the sensor is needed to fulfill basic weather monitoring tasks, it provides sufficient results to be considered a cost-effective solution.

Rain Sensor: The reason why the rain sensor was selected was because of its ability to give versatile outputs of digitized and analogous values to enable it gauge the occurrence and intensity of rainfall. Although it might not be as accurate as professional rain gauges, its flexibility, ability to be easily integrated, and affordability will make it applicable to the real-time weather conditions that the system needs.

3.1.2 Hardware Selection

Cytron Maker Feather S3 Microcontroller: This microcontroller has been chosen since it has built-in Wi-Fi and Bluetooth functions to allow it to easily connect to cloud services. Its integration with MQTT communication protocols as well as the ability to implement lightweight data processing make it work efficiently and reliably. It is a balanced blend of functionality, energy use, and integration simplicity as compared to the less advanced microcontrollers, and so it is suitable in IoT based systems like weather monitoring station.

3.1.3 Scalability

The architecture used in the system would permit the system to **expand in the future** with the ability to add new sensors such as **wind speed or air quality detectors** without making many major changes to the current setup. The services of the **Google Cloud Platform (GCP)** such as **BigQuery** and **Looker Studio** are used to guarantee that the system is able to process the

growing amount of data efficiently, as well as allow **dynamic and real-time visualization**. The scalable design offers flexibility to support the future needs and changing application requirements without the compromise of system performance.

3.1.4 Data Security

The system encompasses several **security protocols** that will guarantee that data is not compromised. **Transport Layer Security (TLS)** has been applied to MQTT to ensure protection against the transmission of data between the IoT devices and the cloud, as well as to ensure data encryption. Firewall rules are implemented with the help of a **Virtual Private Cloud (VPC)**, and the cloud resources are secured against illegal access to the network. Besides, **Identity and Access Management (IAM)** implements the **role-based access control (RBAC)** that means that authorized users and applications have access to system resources only. Even though these security techniques add more complexity to the setup, it is rather much more consistent with the best practices of IoT security and increases the overall rigor and reliability of the system greatly.

3.1.5 Connectivity & Communication

The **Mosquitto MQTT Broker** that is deployed on a **Google Cloud Platform (GCP) Compute Engine** virtual machine is an excellent solution to a reliable and centralized messaging hub of the system. Lightweight **protocol of MQTT** and **PubSubClient** library at microcontroller side and Paho MQTT library at server side protocols allow efficient and **reliable communication** between the hardware layer and the cloud infrastructure. This effective communication system helps to reduce latency and minimise bandwidth usage which are crucial conditions in real-time weather monitoring systems based on IoT.

3.1.6 Cost Effectiveness

Such **low-cost sensors** as the DHT11 sensor and the rain sensor coupled with the **pay-as-you-go** pricing model of Google Cloud Platform will keep the system affordable to a large user base. Mongodb components have been shown to easily ingest real-time data and Big Query provides large-scale data analytics without expending large amounts of money on on-premises infrastructure. The system is able to run without unwarranted spending by using cloud services that are managed and by utilizing affordable hardware. Even though the chosen components and settings have a **better balance of quality**, the trade-off between functionality,

performance, and the overall cost, it can be stated that the advantage lies in more sophisticated sensors and better cloud options which might increase performance.

3.2 Compromises

3.2.1 Design Compromises Between Cost and Performance

The choice of **cost-effective sensors** like the **DHT11** and the **simple rain sensor** was taken deliberately because they are **less precise** than the expensive ones. An example is the **DHT11 which is not as accurate as the DHT22** but is **accurate enough to give real time weather conditions**. The rain sensor, on the same note, is only useful in detecting precipitation, but **not providing quantitative precipitation data** as advanced rain gauges would do. The trade-offs on these designs were necessary in order to have a cheaper real time weather monitoring system at the expense of the essential functionality of the system.

3.2.2 Usability vs. Security Considerations

The use of **advanced security settings**, which is tight **firewall rules** and **encrypted MQTT communication**, made it **more complicated to deploy** and **maintain a system**. Even though such measures are important to guarantee good protection of data, it **needs further configuration** and **technical support**, which can **promote the learning curve of users**. A choice or viewpoint in which it is clear that **security is more than ease of use** depicts the significance of **data integrity and assurance** in the **user of IoT applications**. This tradeoff on design makes sure that the system is **secure and reliable** even though this will **minimally decrease accessibility** with respect to **non-technical users**.

3.2.3 Real-Time Data vs. Long-Term Analysis

The system architecture focuses on the **real-time data corruption** with **MongoDB** and **Big Query** is used to store and **analyze data in the long run**. Even though it is more **efficient** and **scalable in terms of approach**, it **adds some complexity** to the data pipeline. This **trade-off** in design is explained by the fact that such system allows supporting **real-time operation** as well as **analysis of historical data** fulfilling its requirements.

3.2.4 Flexibility vs. Specialization

The choice in favor of **general-purpose components** including **Maker Feather S3** and a **rain sensor** has a tradeoff against **precision and specialization** and leads to **system flexibility**. An example of this is that implementing a specific weather station kit would be easy to integrate, however, it will **decrease the options of tailoring the system** to meet a particular use case.

4. Development Process

4.1 Setting Up Google Cloud Platform Services

Services that are used for the development:

- Compute Engine (VM Instances)
- VPC Network (Firewall Rule)
- IAM & Admin (Service Accounts)
- BigQuery
- Looker Studio

4.1.1 Compute Engine

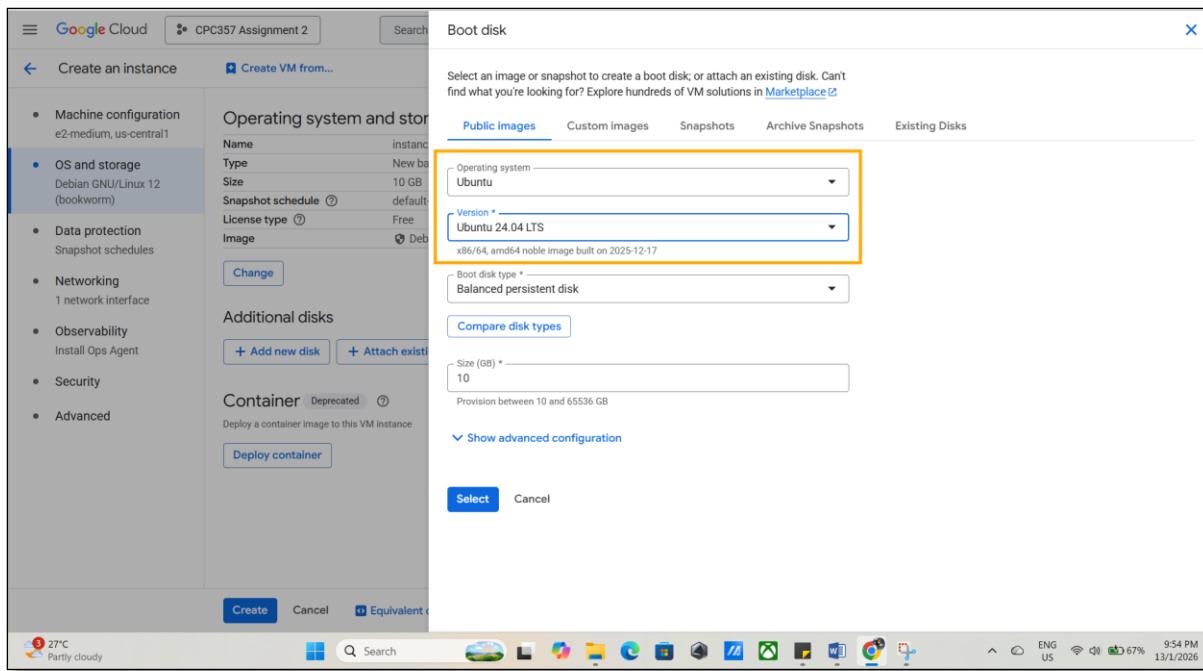
Compute Engine is a scalable compute engine enabling you to establish and operate virtual machines on the Google cloud. We have the option of having ready-to-do and pre-built configurations of VMs to start any machine quickly or we can opt to have custom types of machines and build VMs with the most required processing power and memory depending on our requirement. Click this link to enter into the Google Cloud Console. This will require a few more steps before you will be able to continue the rest of the session if you are working with Compute Engine the first time. Depending on the console, please create a project and enable the Compute Engine API. It is also important to mention that to use the service, you will also have to create a billing account.

Next, go to the navigation menu and select **Compute Engine > Virtual Machines > VM Instances**. Click the **Create Instance** button to create a new virtual machine.

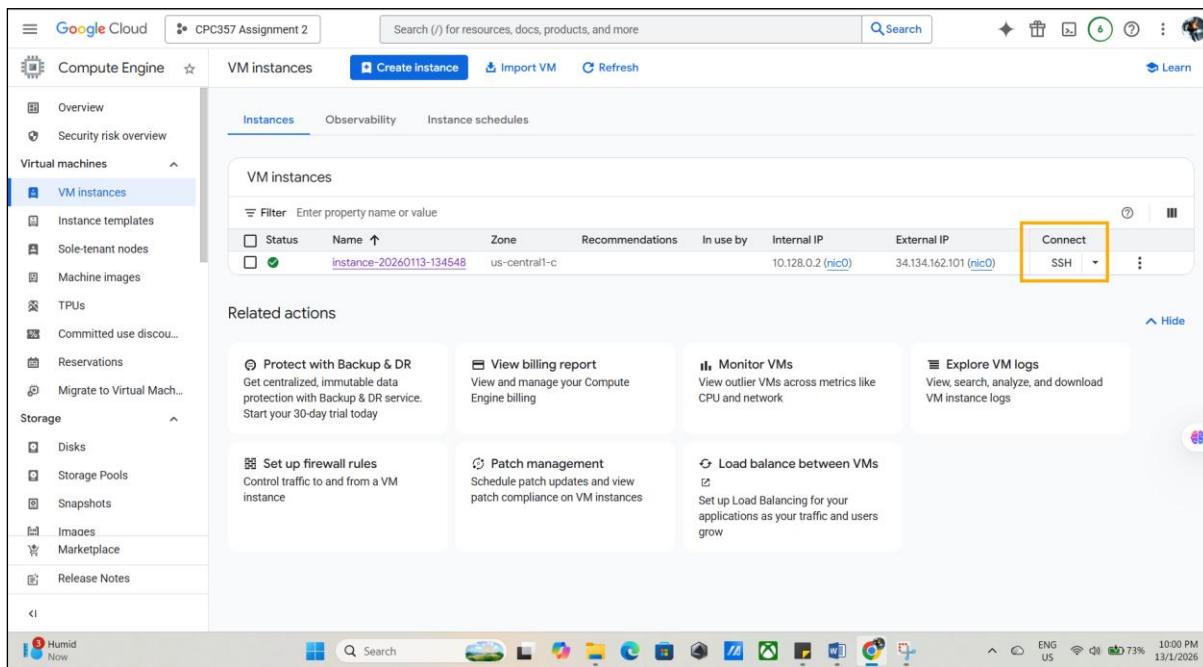
The screenshot shows the Google Cloud Compute Engine VM Instances page. The left sidebar has 'VM instances' selected. At the top, there's a 'Create Instance' button highlighted with a yellow box. Below it, there's a table of VM instances with one entry: 'instance-20260113-134548'. Underneath the table, there's a section titled 'Related actions' with several options like 'Protect with Backup & DR', 'View billing report', 'Monitor VMs', 'Explore VM logs', 'Set up firewall rules', 'Patch management', and 'Load balance between VMs'.

Under **Machine configuration**, you can set the region and zone that runs your virtual machine. Next, go to **OS and storage** in the navigation menu to choose the operating system. Click the Change button and select **Ubuntu** for the operating system and **Ubuntu 24.04 LTS** for the version.

The screenshot shows the 'Create an instance' dialog. On the left, there's a sidebar with various configuration sections: Machine configuration (selected), OS and storage (Debian GNU/Linux 12 (bookworm)), Data protection (Snapshot schedules), Networking (1 network interface), Observability (Install Ops Agent), Security, and Advanced. The main area is titled 'Machine configuration' and shows fields for 'Name' (set to 'instance-20260113-135428'), 'Region' (set to 'us-central1 (Iowa)'), and 'Zone' (set to 'Any'). A note says 'Google will choose a zone on your behalf, maximizing VM obtainability. Zone is permanent.' To the right, there's a 'Monthly estimate' table showing costs for 2 vCPU + 4 GB memory (\$24.46), 10 GB balanced persistent disk (\$1.00), Logging, Monitoring, Snapshot schedule, and a total of \$25.46. Below the table are links for 'Compute Engine pricing' and 'Cloud Operations pricing'.



Click the **Create** button to finish up the instance creation. Upon completion, you should be able to find your newly created VM instance in the list. Next, click the **SSH** button next to the VM instance to open up the SSH-in-browser.



A pop-up window will appear and click the **Authorize** button to give permission for the **SSH-in-browser** to connect to the VM.

```

System load: 0.0      Processes: 113
Usage of /: 33.8% of 8.65GB  Users logged in: 0
Memory usage: 8%      IPv4 address for ens4: 10.128.0.2
Swap usage: 0%

Expanded Security Maintenance for Applications is not enabled.

26 updates can be applied immediately.
25 of these updates are standard security updates.
To see these additional updates run: apt list --upgradable

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

harimessi01@instance-20260113-134540:~$ 

```

After that, run the following commands in the SSH-in-browser to setup the VM:

\$ sudo apt-get update

\$ sudo apt-get upgrade

4.1.2 Firewall

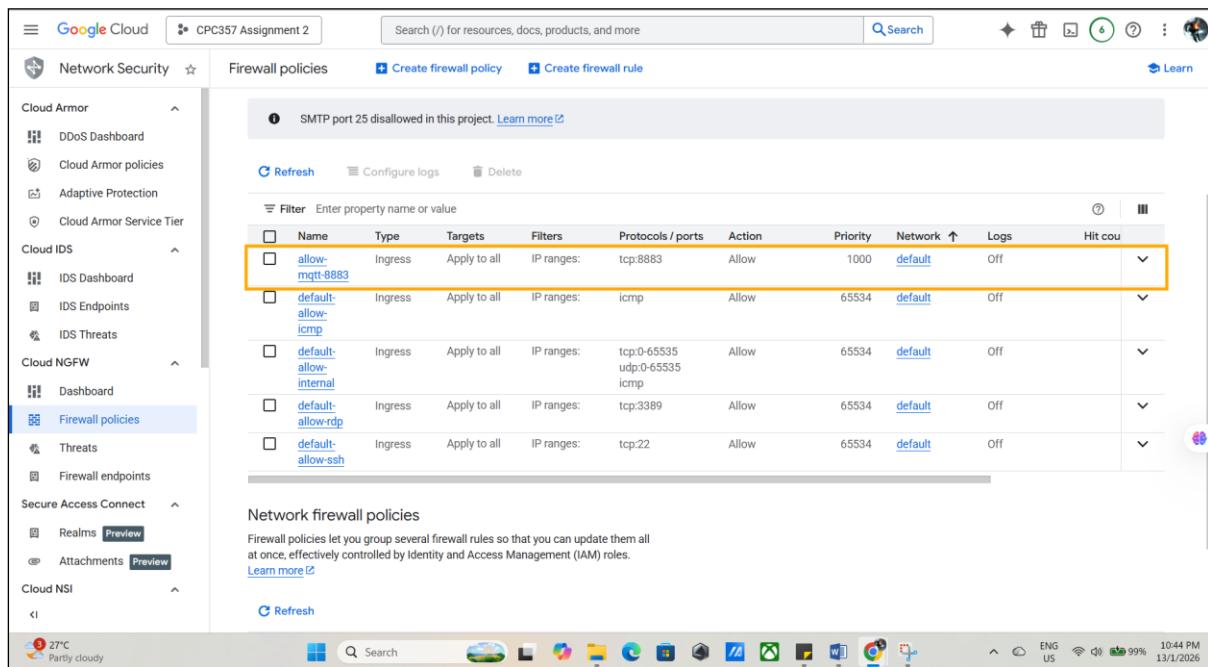
To provide connectivity for your Compute Engine VM instances, you need to configure the Virtual Private Cloud (VPC) network and set up the VPC firewall rules accordingly. To put it simply, VPC firewall rules let you allow or deny traffic to or from VM instances in a VPC network based on port number, tag, or protocol. First, go to the navigation menu and select **VPC Network > Firewall**. Then, click the **Create Firewall Rule** button at the top.

Name	Type	Targets	Filters	Protocols / ports	Action	Priority	Network	Logs	Hit count
allow-mqtt-8883	Ingress	Apply to all	IP ranges:	tcp:8883	Allow	1000	default	Off	
default-allow-icmp	Ingress	Apply to all	IP ranges:	icmp	Allow	65534	default	Off	
default-allow-internal	Ingress	Apply to all	IP ranges:	tcp:0-65535 udp:0-65535 icmp	Allow	65534	default	Off	
default-allow-rdp	Ingress	Apply to all	IP ranges:	tcp:3389	Allow	65534	default	Off	
default-allow-ssh	Ingress	Apply to all	IP ranges:	tcp:22	Allow	65534	default	Off	

Configure the VPC firewall rule according to the specifications below:

- Name: allow-mqtt-8883
- Description: Allow traffic on port 8883
- Logs: off
- Network: default
- Priority: 1000
- Direction of traffic: Ingress
- Action on match: Allow
- Targets: All instances in the network
- Source filter: IPv4 ranges
- Source IPv4 ranges: 0.0.0.0/0
- Second source filter: None
- Destination filter: None
- Protocols and ports: Specified protocols and ports (TCP Port 8883)

Once you're done, click on the Create button to save and create the VPC firewall rule and upon completion, the newly created VPC firewall rule will be displayed in the list.



The screenshot shows the Google Cloud Network Security Firewall policies interface. On the left, there's a navigation sidebar with sections like Cloud Armor, Cloud IDS, Cloud NGFW, Secure Access Connect, and Cloud NSI. Under Firewall policies, there are options for Threats and Firewall endpoints. The main area displays a table of firewall rules. One rule, 'allow-mqtt-8883', is highlighted with a yellow border. The table columns include Name, Type, Targets, Filters, Protocols / ports, Action, Priority, Network, Logs, and Hit count. Other visible rules include 'default-allow-icmp', 'default-allow-internal', 'default-allow-rdp', and 'default-allow-ssh'. Below the table, there's a section titled 'Network firewall policies' with a note about grouping rules for IAM roles.

Name	Type	Targets	Filters	Protocols / ports	Action	Priority	Network	Logs	Hit cou
allow-mqtt-8883	Ingress	Apply to all	IP ranges:	tcp:8883	Allow	1000	default	Off	
default-allow-icmp	Ingress	Apply to all	IP ranges:	icmp	Allow	65534	default	Off	
default-allow-internal	Ingress	Apply to all	IP ranges:	tcp:0-65535 udp:0-65535 icmp	Allow	65534	default	Off	
default-allow-rdp	Ingress	Apply to all	IP ranges:	tcp:3389	Allow	65534	default	Off	
default-allow-ssh	Ingress	Apply to all	IP ranges:	tcp:22	Allow	65534	default	Off	

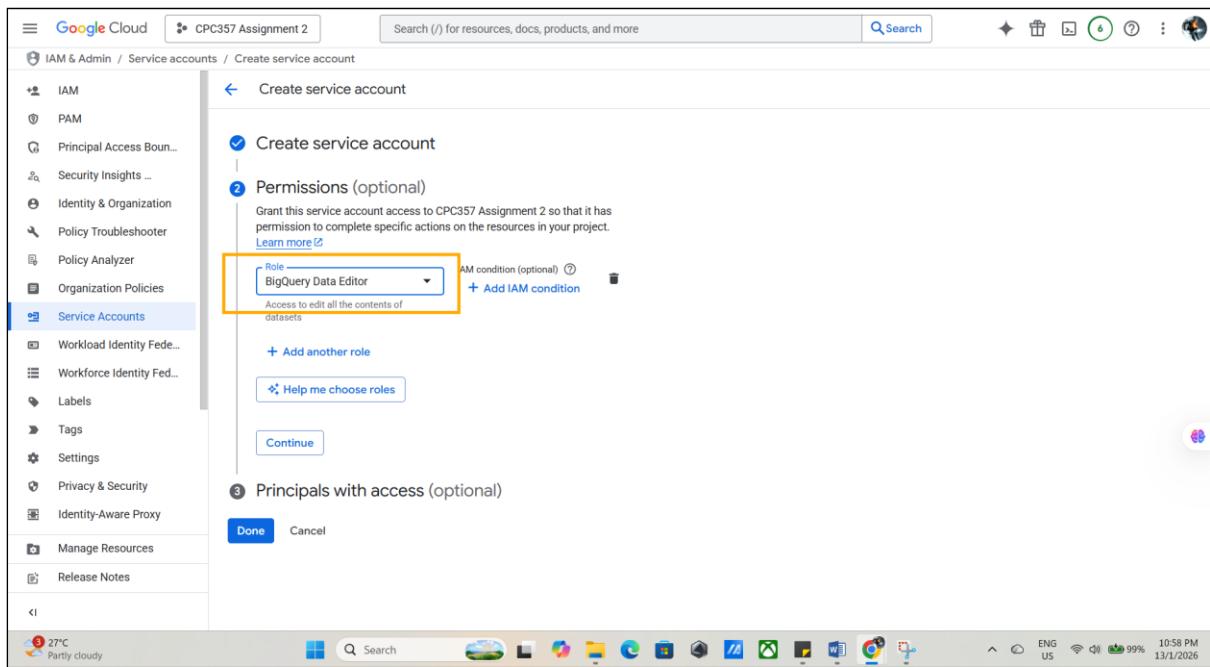
4.1.3 Service Account

A service account represents a Google Cloud service identity, such as code running on Compute Engine VMs, App Engine apps, or systems running outside Google. Organization policies can be used to secure service accounts and block risky service account features, such as automatic IAM Grants, key creation/upload, or the creation of service accounts entirely. Go to the navigation menu and select **IAM & Admin > Service Accounts**. Then, click the **Create Service Account** button at the top.

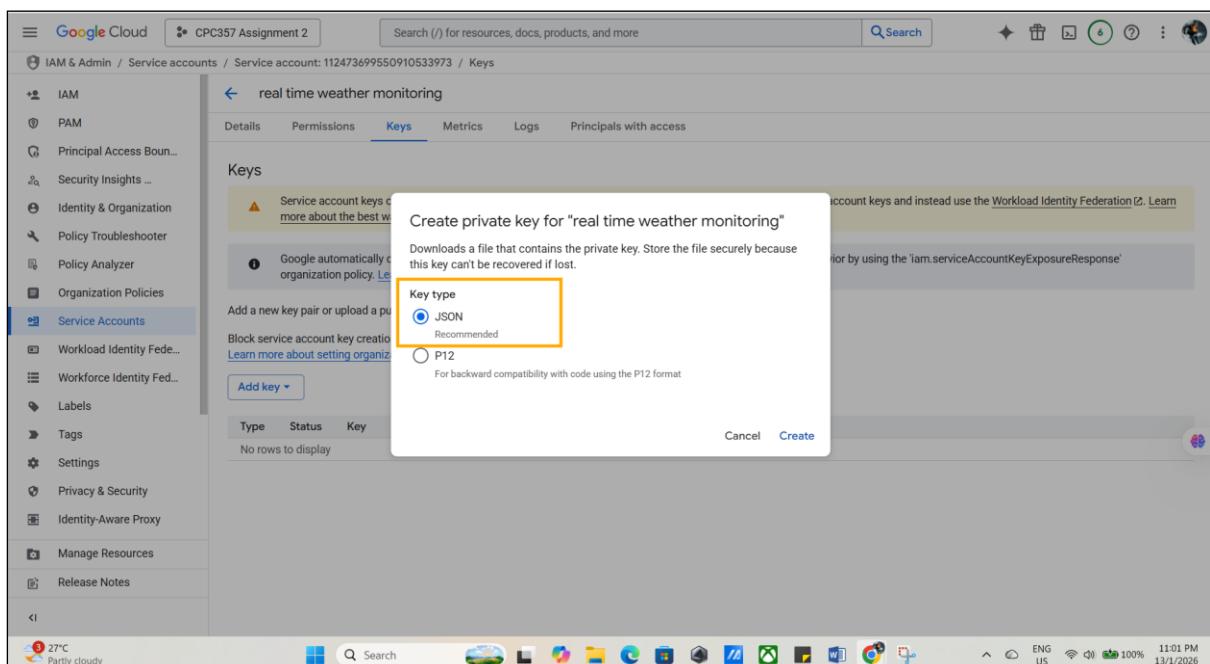
The screenshot shows the Google Cloud IAM & Admin / Service accounts page. The left sidebar is collapsed. The main area displays a table of service accounts. The first account is 'Compute Engine default service account' (Email: 54082342092-compute@developer.gserviceaccount.com), which is Enabled, has 'Compute Engine' as its name, and 'No keys' as its key status. The second account is 'real time weather monitoring' (Email: real-time-weather-monitoring@cpc357-assignment-2-484213.iam.gserviceaccount.com), also Enabled, with 'real time weather monitoring' as its name and 'No keys' as its key status. A blue box highlights the '+ Create service account' button at the top right of the table header.

Under **Service account details**, give your account a name and click **Create and Continue**. Then, grant a role to your service account which is **BigQuery Data Editor** (access to edit all the contents of datasets). After that, create the service account by clicking **Done**.

The screenshot shows the 'Create service account' dialog box. The left sidebar is collapsed. The dialog is divided into three numbered sections: 1. Create service account, 2. Permissions (optional), and 3. Principals with access (optional). Step 1 is highlighted with a yellow box. In step 1, the 'Service account name' field contains 'weather monitoring system'. Below it, the 'Display name for this service account' field is empty. The 'Service account ID' field shows 'weather-monitoring-system'. The 'Email address' field shows 'weather-monitoring-system@cpc357-assignment-2-484213.iam.gserviceaccount.com'. The 'Service account description' field is empty. At the bottom of step 1 is a 'Create and continue' button. At the bottom of the dialog are 'Done' and 'Cancel' buttons. The status bar at the bottom indicates it's 10:57 PM on 13/1/2026.

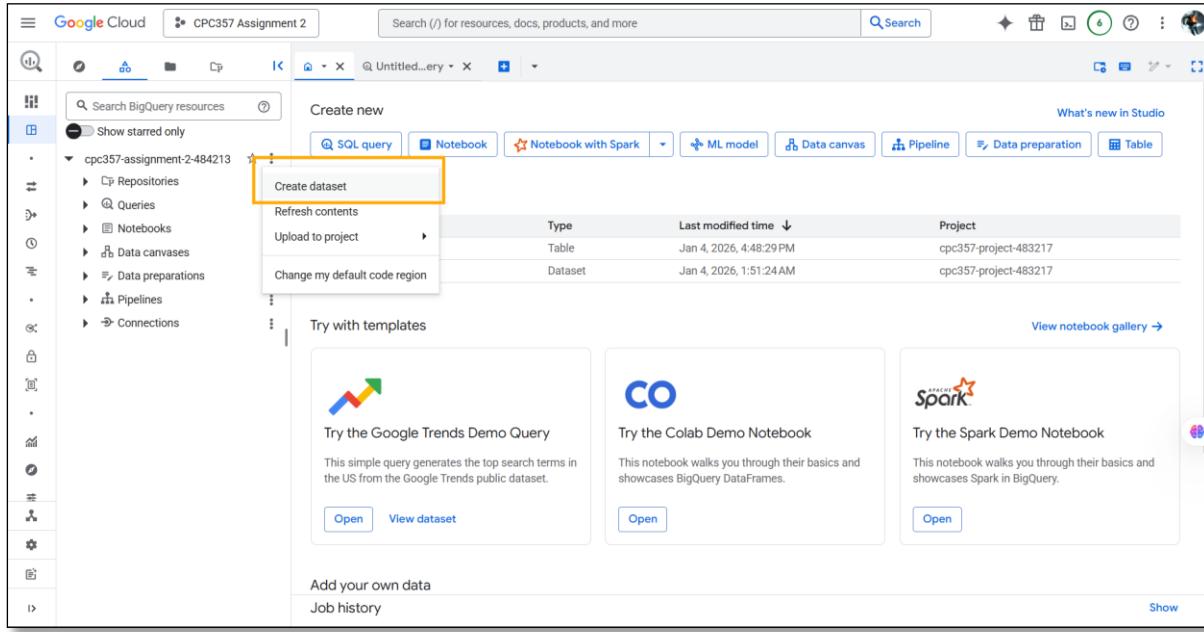


After creating the service account, the account will be displayed in the list. Click the account and navigate to **Keys** and click **Add Key** > **Create New Key**. Then, in the pop-up window, choose **JSON** as the key type and click **Create**. After clicking, a JSON file containing the service account's key will be downloaded automatically. Keep the file safe because we will use that key in our python script later.

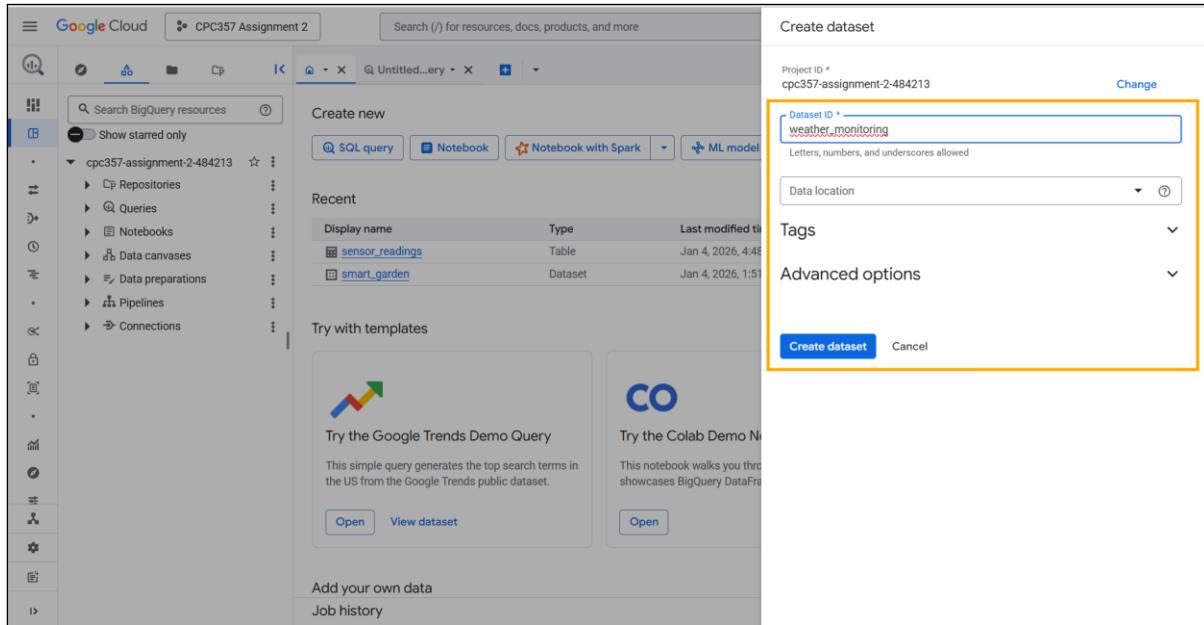


4.1.4 BigQuery

BigQuery is a serverless, cost-effective, and multicloud data warehouse designed to help you turn big data into valuable business insights. Navigate to BigQuery and click **Create dataset** under your project-id. Then, give enter your **Dataset ID** and specify the location type that you want to use to store your data. After that, click **Create**.



The screenshot shows the Google Cloud BigQuery interface. On the left, there's a sidebar with various project resources like Repositories, Queries, Notebooks, Data canvases, Data preparations, Pipelines, and Connections. The main area has a search bar at the top. Below it, there's a 'Create new' section with several options: SQL query, Notebook, Notebook with Spark, ML model, Data canvas, Pipeline, Data preparation, and Table. A dropdown menu is open over the 'Table' option, showing 'Create dataset'. This menu also includes 'Refresh contents', 'Upload to project', and 'Change my default code region'. To the right of this, there's a table showing existing datasets: 'Table' (Jan 4, 2026, 4:48:29 PM) and 'Dataset' (Jan 4, 2026, 1:51:24 AM). Below the table, there are three demo templates: 'Try the Google Trends Demo Query', 'Try the Colab Demo Notebook', and 'Try the Spark Demo Notebook', each with an 'Open' button. At the bottom, there are links for 'Add your own data' and 'Job history'.



This screenshot shows the 'Create dataset' dialog box. It has a 'Project ID' field set to 'cpc357-assignment-2-484213' with a 'Change' link. The 'Dataset ID' field is highlighted with a yellow box and contains the value 'weather_monitoring'. Below it, there's a note: 'Letters, numbers, and underscores allowed'. There's also a 'Data location' dropdown. The dialog box has sections for 'Tags' and 'Advanced options'. At the bottom right are 'Create dataset' and 'Cancel' buttons.

Upon completion, **Create table** under the newly created dataset and give table a name. After that, create four fields as below table under the **Schema** section:

Field name	Type	Mode
timestamp	TIMESTAMP	REQUIRED
temperature	FLOAT	NULLABLE
humidity	FLOAT	NULLABLE
rain	INTEGER	NULLABLE

The screenshot shows the Google Cloud BigQuery interface. On the left, there's a sidebar with various project and dataset options. In the center, a search bar and a 'Create new' section are visible. Below that is a 'Recent' section listing datasets and tables. To the right, there are promotional cards for 'Google Trends Demo Query', 'Try the Colab Demo Notebook', and 'Try the Spark Demo Notebook'. A context menu is open over the 'weather_monitoring' dataset, with the 'Create table' option highlighted.

The screenshot shows the 'Create table' dialog box. It has sections for 'Source' (set to 'Empty table'), 'Destination' (Project: cpc37-assignment-2-484213, Dataset: weather_monitoring, Table: sensor_readings), and 'Schema'. The schema defines two columns: 'timestamp' (Type: TIMESTAMP, Mode: REQUIRED) and 'temperature' (Type: FLOAT, Mode: NULLABLE). There are also buttons for 'Create table' and 'Cancel' at the bottom.

Then click **Create Table**. Now you can see the created table with the column names and their types.

The screenshot shows the Google Cloud BigQuery interface. The left sidebar lists projects and datasets, with 'cpc357-assignment-2-484213' selected. Under 'weather_monitoring', the 'sensor_readings' table is selected. The main pane displays the table schema:

Field name	Type	Mode	Description	Key	Collation	Default Value	Policy Tags	Data Policies
timestamp	TIMESTAMP	REQUIRED	-	-	-	-	-	-
temperature	FLOAT	NULLABLE	-	-	-	-	-	-
humidity	FLOAT	NULLABLE	-	-	-	-	-	-
rain	INTEGER	NULLABLE	-	-	-	-	-	-

4.2 Set Up the Hardware

4.2.1 Arduino Code

Copy the code from the GitHub Repo:

https://github.com/harinethiran1306/Real-Time-Weather-Monitoring-System/blob/fa80ae496471f8addeb1b6d2ea7fc827dace89d7/CPC357_Assignment2_Arduino_Code.ino

Paste it in your Arduino IDE. Before uploading into your microcontroller, make sure you download and enable PubSubClient library in Arduino IDE.

4.3 Set Up Environment in Google Virtual Machine

Run all the following commands before creating the python script:

- **sudo apt-get install mosquito**
- **sudo apt-get install mosquitto-clients**
- **sudo apt install python3-pip**
- **\$ pip install paho-mqtt**
- **sudo apt-get install -y mongodb**
- **pip install pymongo**
- **pip3 install pymongo google-cloud-bigquery paho-mqtt**

4.4 Create python script & Run it

Use the python script from GitHub Repository Link:

<https://github.com/harinethiran1306/Real-Time-Weather-Monitoring-System/blob/fa80ae496471f8addeb1b6d2ea7fc827dace89d7/mongo-by.py>

Simple follow the below command to create and run the script.

- **nano mongo-bq.py**

Then copy the code and paste into the nano editor. After that, click **Ctrl+O** to save and **Ctrl+X** to exit the nano editor. You have to include your service account key in the same directory too. Copy the content from the JSON file, create a **service-account-key.json** file using the same command that used to create the python script, and paste it into the file. Save and close the file then run the below command to execute python script.

- `python3 mongo-bq.py`

You will see the below output indicating the sensor readings are successfully inserted into BigQuery table.

Now navigate to the BigQuery table and click the **Export > Explore with Looker Studio**. You will be redirect to Looker Studio where you can create charts according to your preference.

Row	timestamp	temperature	humidity	rain
55	2026-01-13 18:02:41.858894 UTC	30.8	79.0	1
56	2026-01-13 18:02:46.898725 UTC	30.8	79.0	1
57	2026-01-13 18:02:51.98389 UTC	30.8	79.0	1
58	2026-01-13 18:02:57.318596 UTC	30.8	79.0	1
59	2026-01-13 18:03:01.968669 UTC	30.8	79.0	1
60	2026-01-13 18:03:07.028542 UTC	30.8	79.0	1
61	2026-01-13 18:03:12.029490 UTC	30.8	79.0	1
62	2026-01-13 18:03:17.632491 UTC	30.8	79.0	1
63	2026-01-13 18:03:22.078356 UTC	30.8	79.0	1
64	2026-01-13 18:03:27.118913 UTC	30.8	79.0	1
65	2026-01-13 18:03:32.107974 UTC	30.8	79.0	1
66	2026-01-13 18:03:37.465829 UTC	30.8	79.0	1
67	2026-01-13 18:03:42.230040 UTC	30.8	79.0	1

5. Security Consideration

5.1 Security Measure Implemented

5.1.1 Transport Layer Security (TLS) for MQTT Communication

All messages between the **microcontroller** and the MQTT broker go through encryption using **TLS over port 8883** to guarantee **data confidentiality** and prevent the occurrence of **eavesdropping and man-in-the-middle attacks** during a transmission. Such an encryption means that data packets which have been intercepted **cannot be deciphered** or **modified** correctly without authorization decryption keys. Besides that, TLS certificates also approach the client and server authentication, creating a **reliable connection** between them and avoiding communication with **unauthorized servers**.

5.1.2 Firewall Rules in Virtual Private Cloud (VPC)

Firewall policies in the **Virtual Private Cloud (VPC)** are set to **control the traffic in and out** so as to assure that only the **approved devices and IP ranges** can target the **MQTT broker** and other system resources. In order to be less exposed to vulnerable points, there is a list of ports that are opened:

- **Only encrypted MQTT communication is enabled in Port 8883.**
- **The SSH (port 22) is only limited to the particular IP addresses.**

Intense access control has a significant effect on **minimizing the attack surface** of the system and **eliminating non-authorized access** to cloud elements, which can guarantee **network access security**.

5.1.3 Identify & Access Management (IAM)

The system will take advantage of the **Google Cloud IAM** contributory to **execute role-based access control (RBAC)**, can grant each service and user the bare **minimum permissions** to execute their tasks. The **principle of least privilege** as it is applied minimizes the chance of **unintentional or malicious abuse** of resources.

Examples include:

- **Designation of the role of BigQuery Admin to service accounts that solely handle the management of the data and data queries.**
- **Limiting the end-user accounts to read-only access of the Looker Studio dashboards, excluding any possibilities of altering or deleting the data.**

The IAM policies are continuously revised in order to keep in pace with the changing requirements of the systems, thus **ensuring enhanced security** in the long run.

5.1.4 Data Encryption in Transit & at Rest

TLS has been deployed in **MQTT** communications and **HTTPS** in accessing dashboards, which helps in encrypting data during these interactions. Such encryption discourages the revelation of delicate information as data transfers between parts of the systems. Also, every piece of data in both **BigQuery** and the **MongoDB** is encrypted in place with the default **encryption of AES-256** that is provided by the Google Cloud. This is so that in case of any physical attacks on the storage devices, there will be no security risk on the data and the data will not be readable in an unipolar manner without the decryption keys. The encryption features of Google Cloud are up to industry standards that give users the system of **trust and reliability**.

5.1.5 Authentication & Authorization

There is authentication of the microcontroller, MQTT broker and cloud services using secure **API keys** and **OAuth tokens**. These tokens allow users to be sure that only **true devices** and **users** can log in to the endpoints of the system. Also, API keys are regularly changed and the lifetime of tokens is limited to ensure that they are not abused in case of a compromise. Each

of the **API endpoints** is set to deny unauthenticated requests to give a structure allowing the **blocking of unauthorized admittance**.

5.2 Strategies for Data Protection & Access Control

5.2.1 Monitoring & Alerts

Real-time monitoring is established using **Google Cloud Monitoring** that monitors the activity of the system and detects abnormal patterns. Alerts are set to inform the administrators that a **possible breach** has occurred, like a request to access the information without an authorization, or an irregularity in the data flows. Logs of **Cloud Audit Logs** provide administrators with an opportunity to detect and act upon suspicious activity in a short period of time. Such constant tracking guarantees that there are **timely interventions** to reduce risks and guarantee security of data.

5.2.2 Backup & Disaster Recovery

A **powerful backup strategy** is used in the system, where key data is saved on safe places on a regular basis. This is to guarantee fast restoration of data in case they are breached or the system fails. There is also a **disaster recovery plan** to recover services in **minutes** so that there is minimum downtimes. Such measures ensure the system is not affected by **data loss** and offers operational durability in case of **unforeseen incidents**.

5.2.3 Access Control Policies

The system implements **granular access control** to provide that only **authorized devices** and authorized users can access certain resources. **Google Cloud Identity and Access Management (IAM)** is used to implement **role-based access control (RBAC)** to allocate roles with minimal permissions needed to do tasks. As an example, the role of the **BigQuery Admin** provides access to the data stored information to be managed, whereas **end-user accounts** are allowed to view the visualizations in **Looker Studio** only without writing something. These policies play a vital role in ensuring that there is **integrity in the system** and unauthorized access is avoided.

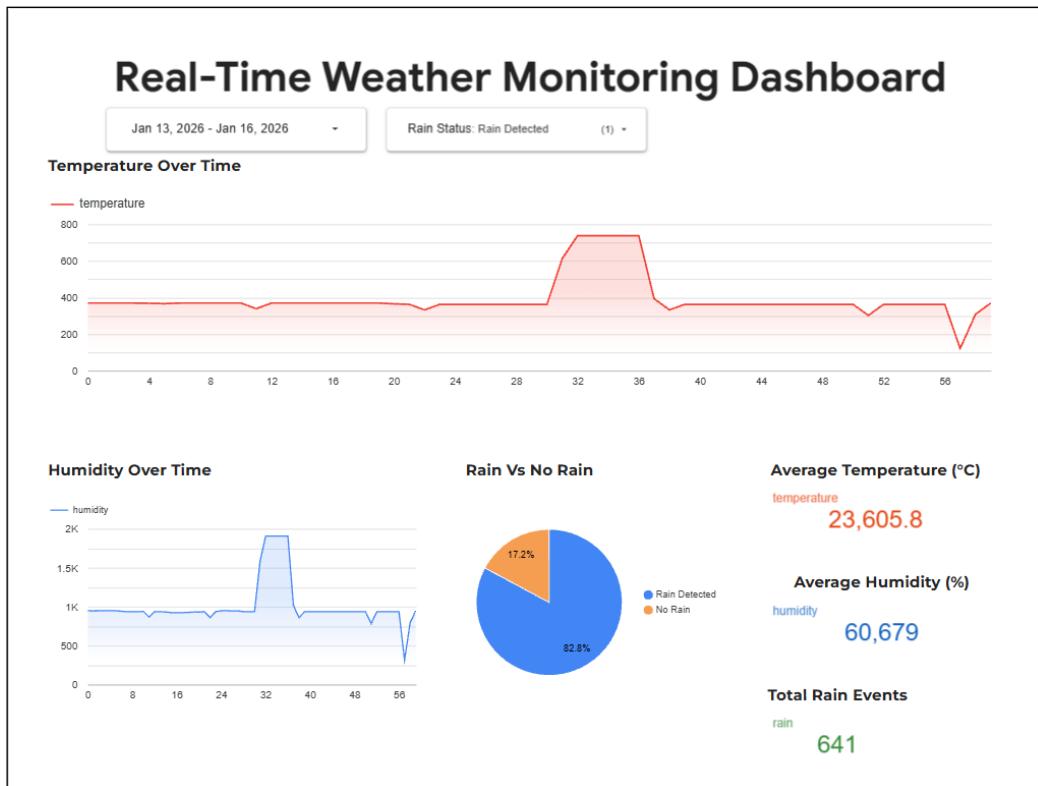
5.2.4 Regular Key Rotation

To reduce the chances of breached credentials, the system will carry out **frequent key rotation** of all **encryption keys** and **API keys**. This habit also means that in case a key is left exposed, it has a restricted time during which it would be used. Automation of this process by the system will lighten the load of the administration and improve the overall **security posture**. This is also an offensive way of ensuring **data confidentiality** and reducing the exposure to possible threats.

5.2.5 Penetration Testing & Vulnerability

Periodic **penetration testing** of the system is done to discover and eliminate vulnerabilities that may be exploited by the attackers. Automated **vulnerability scans** are also done to identify **misconfigurations**, old libraries or open endpoints. These weaknesses can be mitigated in advance to guarantee the **integrity of the data** and lower the chances of successful cyberattacks.

6. Appendix



- GitHub Repository Link: <https://github.com/harinethiran1306/Real-Time-Weather-Monitoring-System.git>